

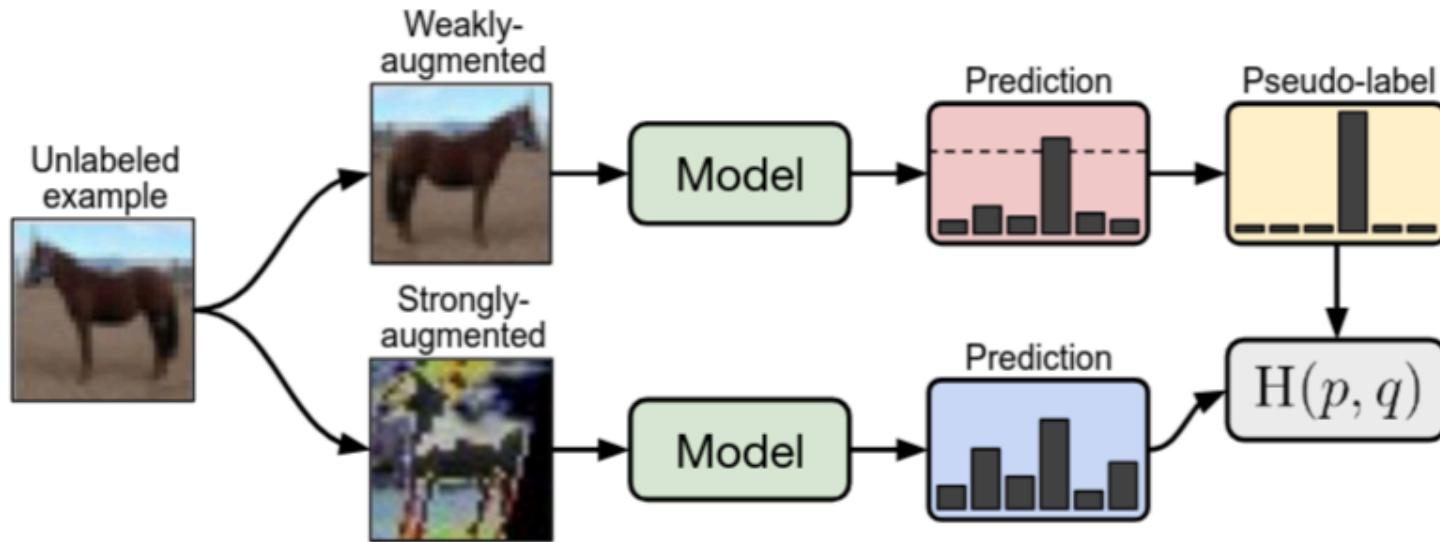
Deep Learning - Assignment 3

Semi Supervised Learning com o algorítmo FixMatch.

Guilherme Castilho, Pedro Tokar e Vitor do Nascimento

Fluxo de trabalho

- Plataformas e ferramentas utilizadas: Google Colab, Kaggle e VS Code;
- Criação das funções necessárias para a experimentação;
- Experimentação com diferentes números de samples por classe;
- Experimentação com cut-out;
- Experimentação com amostras representativas e não representativas;



Todos os testes foram executados com os hiperparâmetros e schedulers listados no paper:

- Batch size 64;
- Learning rate: 0.003;
- Weight decay: 0.0005;
- Parâmetros da Loss: $\tau = 0.95$, $\lambda_u = 1$;
- Proporção do batch sem labels: $\mu = 7$;
- Scheduler com função $\cos\left(\frac{7\pi k}{16K}\right)$;

Implementação

```
#Labeled dataset
def __getitem__(self, idx):
    real_idx = self.indexes[idx]
    image, label = self.dataset[real_idx]
    image = image.to(torch.float).to(self.device)

    image = self.weak_augmentations(image)
    if NORMALIZE:
        image = self.normalize_img(image)

    return image, label

#Unlabeled dataset
def __getitem__(self, idx):
    real_idx = self.indexes[idx]
    image, label = self.dataset[real_idx]
    image = image.to(torch.uint8).to(self.device)

    weak_image = self.weak_augmentations(image).to(torch.float)
    strong_image = self.strong_augmentations(image).to(torch.float)
    if NORMALIZE:
        weak_image = self.normalize_img(weak_image)
        strong_image = self.normalize_img(strong_image)

    return weak_image, strong_image
```

```
def get_split_dataset(dataset, n_classes, n_samples, ctaugment = True, device = "cpu"):
    labeled_indexes = []

    random_indexes = torch.randperm(len(dataset))

    total_indexes = n_classes * n_samples
    frequencies = defaultdict(lambda: 0)
    curr = 0

    while len(labeled_indexes) < total_indexes:
        if curr == len(dataset):
            raise RuntimeError("Não foi possível fazer split do dataset")

        label = dataset[random_indexes[curr]][1]
        if frequencies[label] < n_samples:
            labeled_indexes.append(random_indexes[curr].item())
            frequencies[label] += 1
        curr += 1

    unlabeled_indexes = list(set(range(len(dataset))) - set(labeled_indexes))

    labeled_dataset = LabeledDataset(dataset, labeled_indexes, device)
    unlabeled_dataset = UnlabeledDataset(dataset, unlabeled_indexes, ctaugment, device)
    return labeled_dataset, unlabeled_dataset
```

```
class FixMatchLoss(nn.Module):
    def __init__(self, threshold, weight):
        super().__init__()
        self.threshold = threshold
        self.weight = weight
        self.cross_entropy = nn.CrossEntropyLoss(reduction = "sum")
        self.softmax = nn.Softmax(dim = 1)

    def forward(self, labeled_predictions, labeled_truth, unlabeled_weak_predictions, unlabeled_strong_predictions):
        l_s = self.cross_entropy(labeled_predictions, labeled_truth)/labeled_truth.shape[0]

        with torch.no_grad():
            unlabeled_weak_predictions = self.softmax(unlabeled_weak_predictions)
            mask = unlabeled_weak_predictions.max(dim = 1)[0] > self.threshold
            if sum(mask) == 0:
                return l_s, 0

            l_u = self.cross_entropy(
                unlabeled_strong_predictions[mask],
                unlabeled_weak_predictions[mask].argmax(dim = 1)
            )/unlabeled_weak_predictions.shape[0]

        return l_s + self.weight * l_u, sum(mask)
```

```
zip_loaders = zip(cycle(self._labeled_dataloader), self._unlabeled_dataloader)
progress_bar = tqdm(zip_loaders, desc = f"Epoch {epoch + 1} | Training", total = len(self._unlabeled_dataloader))

for i, (labeled_data, unlabeled_data) in enumerate(progress_bar):
    labeled_images, labeled_labels = labeled_data
    unlabeled_weak, unlabeled_strong = unlabeled_data
    #[...]
    self._optimizer.zero_grad()

    labeled_predictions = self._model(labeled_images)
    unlabeled_weak_predictions = self._model(unlabeled_weak)
    unlabeled_strong_predictions = self._model(unlabeled_strong)

    loss, confident = self._criterion(labeled_predictions,
                                        labeled_labels,
                                        unlabeled_weak_predictions,
                                        unlabeled_strong_predictions)
    loss.backward()
    self._optimizer.step()

    running_loss += loss.item()
    progress_bar.set_postfix(loss = loss.item(), lr = self._optimizer.param_groups[0]["lr"])
    confidents += int(confident)

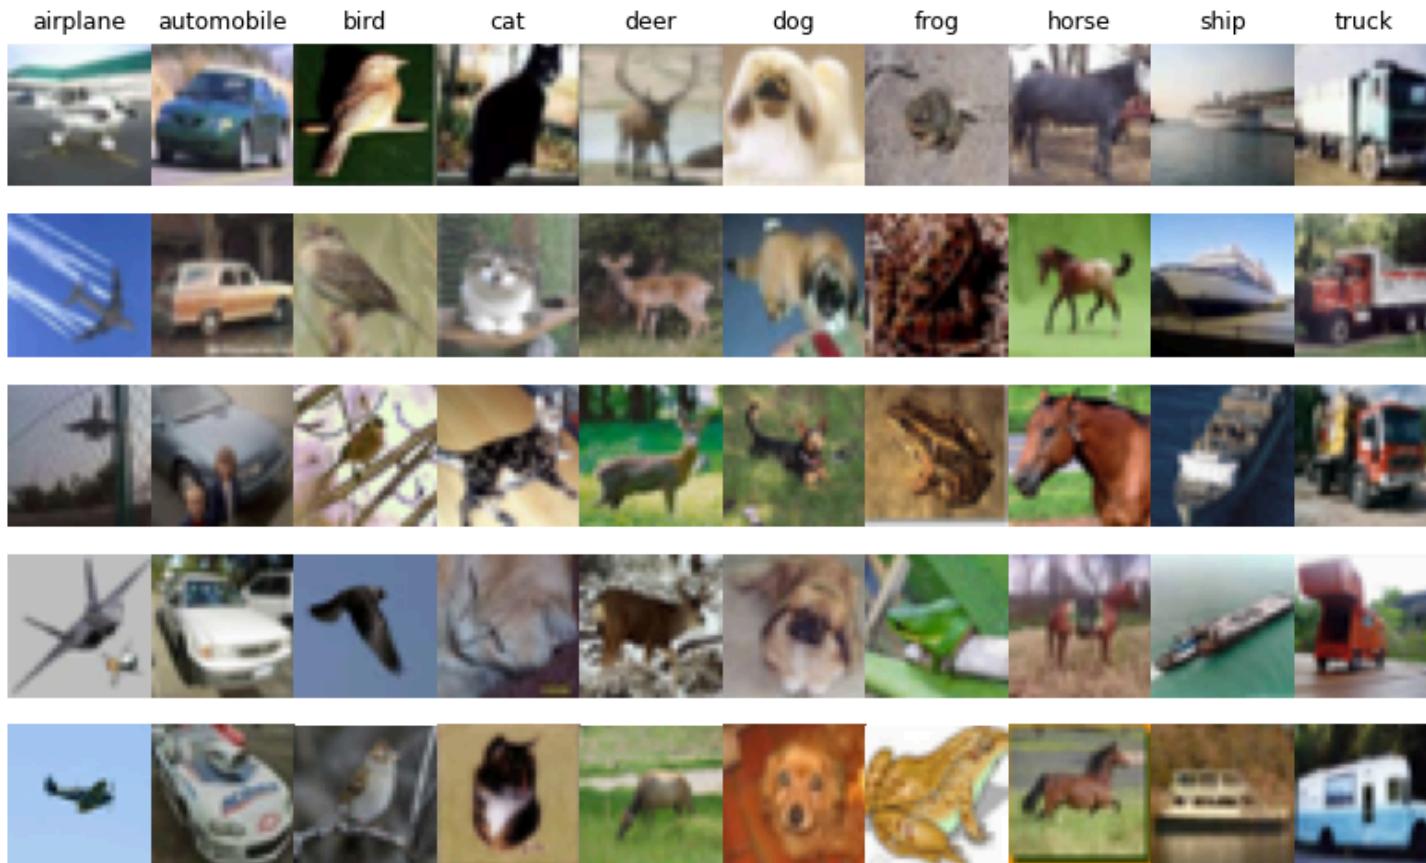
epoch_loss = running_loss / len(self._unlabeled_dataloader)
percentage_confiance = 100*confidents/len(self._unlabeled_dataloader.dataset)
```

```
optimizer = optim.SGD(
    model.parameters(),
    lr = learning_rate,
    momentum = 0.9,
    nesterov = True,
    weight_decay = 0.0005
)
scheduler = optim.lr_scheduler.LambdaLR(
    optimizer,
    lambda epoch: cos((7*pi*epoch) / (16*epochs))
)
loss_function = FixMatchLoss(threshold, lambda_u)

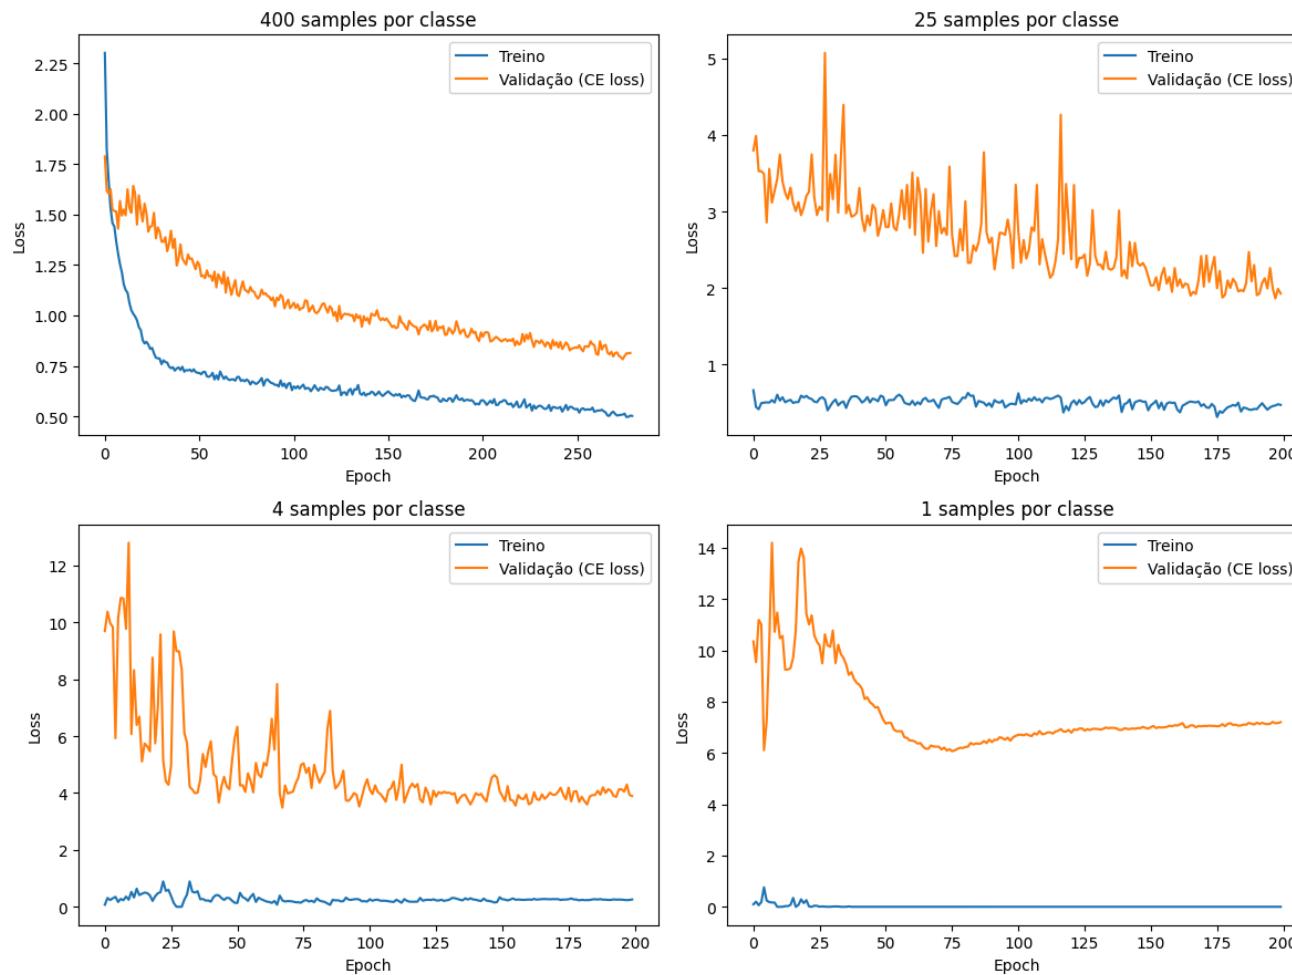
trainer = Trainer(
    model,
    labeled_dataset,
    unlabeled_dataset,
    loss_function,
    optimizer,
    val_dataset = test_dataset,
    device = DEVICE
)
```

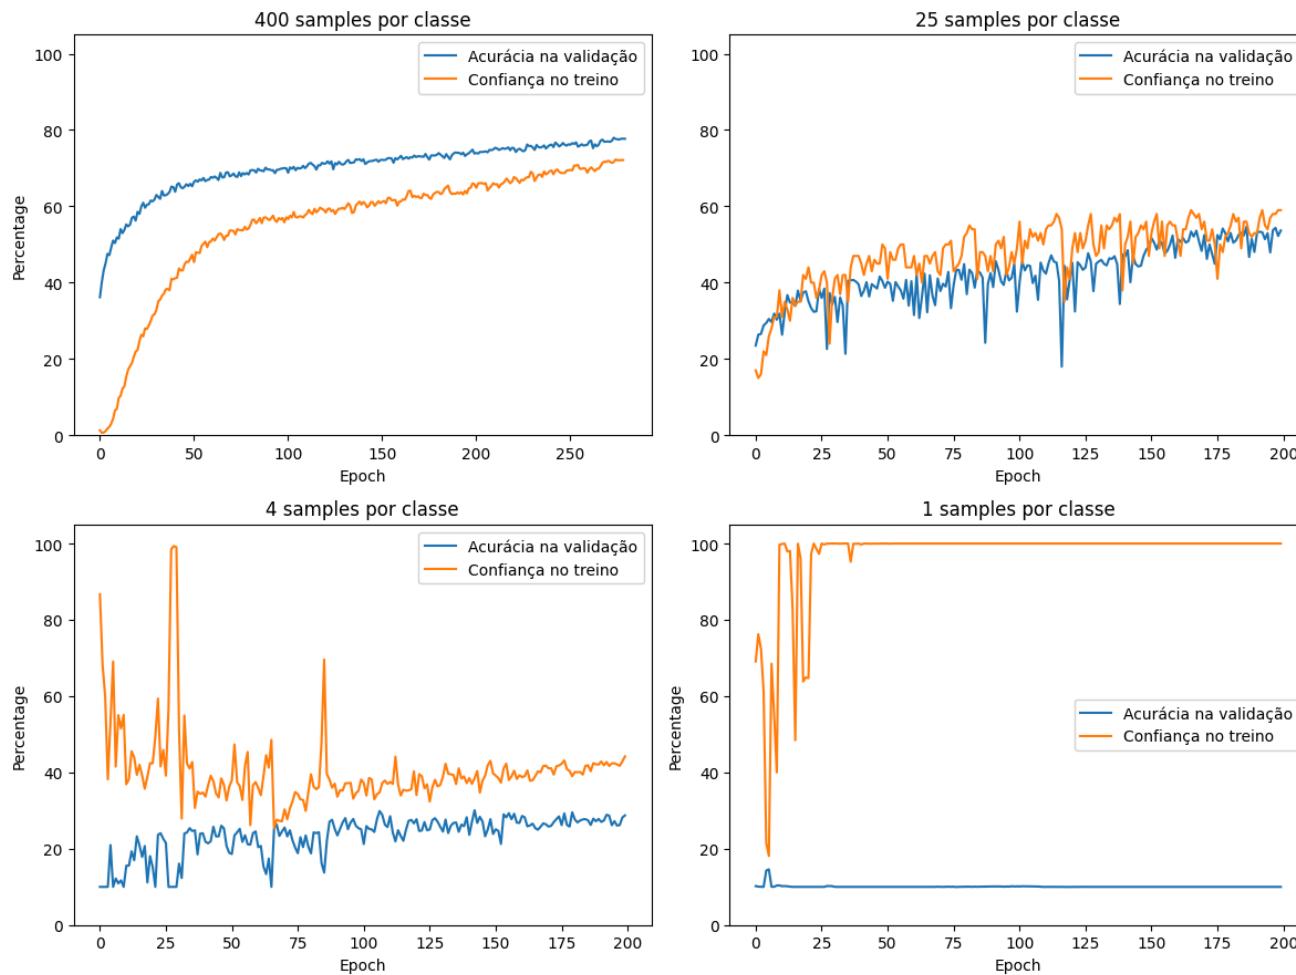
Imagenes

Labeled Dataset Images



Experimento: número de samples





Augmentation

- AutoAugment é um método para aprender uma política de aumento de dados que resulta em alta acurácia no conjunto de validação;
- Política de aumento: sequência de tuplas (transformação, magnitude do parâmetro);
- No AutoAugment a política é aprendida com supervisão;
- RandAugment: Resolve a questão do treinamento supervisionado por meio de amostragem aleatória das transformações, mas exige ajuste nos hiperparâmetros;
- CTAugment: Assim como o RandAugment, também amostra transformações uniformemente ao acaso, mas infere dinamicamente as magnitudes para cada transformação durante o processo de treinamento.

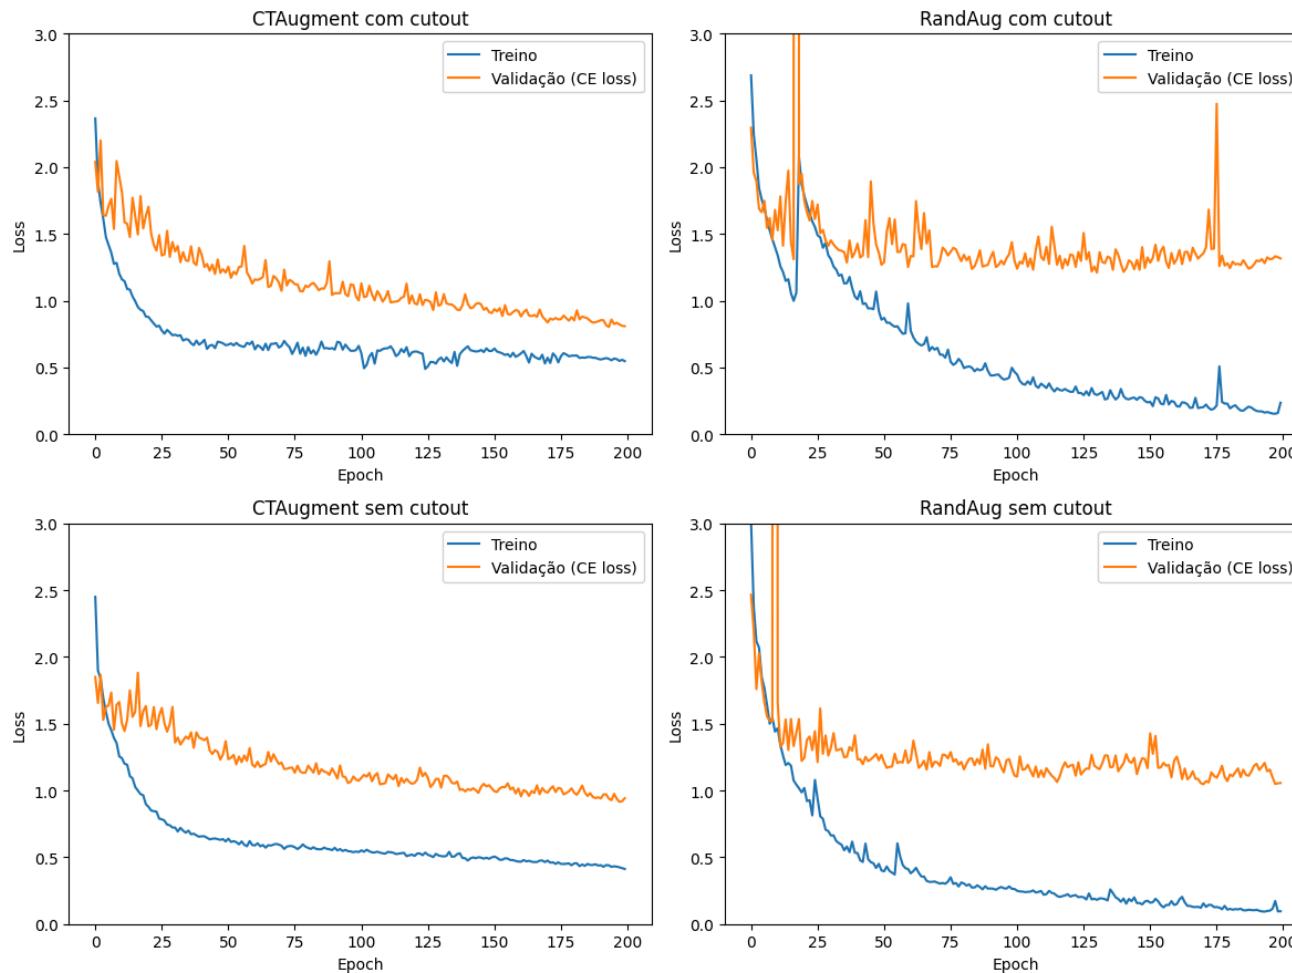
Transformation	Description	Parameter	Range
Autocontrast	Maximizes the image contrast by setting the darkest (lightest) pixel to black (white).		
Brightness	Adjusts the brightness of the image. $B = 0$ returns a black image, $B = 1$ returns the original image.	B	[0.05, 0.95]
Color	Adjusts the color balance of the image like in a TV. $C = 0$ returns a black & white image, $C = 1$ returns the original image.	C	[0.05, 0.95]
Contrast	Controls the contrast of the image. A $C = 0$ returns a gray image, $C = 1$ returns the original image.	C	[0.05, 0.95]
Equalize	Equalizes the image histogram.		
Identity	Returns the original image.		
Posterize	Reduces each pixel to B bits.	B	[4, 8]
Rotate	Rotates the image by θ degrees.	θ	[-30, 30]
Sharpness	Adjusts the sharpness of the image, where $S = 0$ returns a blurred image, and $S = 1$ returns the original image.	S	[0.05, 0.95]
Shear_x	Shears the image along the horizontal axis with rate R .	R	[-0.3, 0.3]
Shear_y	Shears the image along the vertical axis with rate R .	R	[-0.3, 0.3]
Solarize	Inverts all pixels above a threshold value of T .	T	[0, 1]
Translate_x	Translates the image horizontally by $(\lambda \times \text{image width})$ pixels.	λ	[-0.3, 0.3]
Translate_y	Translates the image vertically by $(\lambda \times \text{image height})$ pixels.	λ	[-0.3, 0.3]

Transformation	Description	Parameter	Range
Autocontrast	Maximizes the image contrast by setting the darkest (lightest) pixel to black (white), and then blends with the original image with blending ratio λ .	λ	[0, 1]
Brightness	Adjusts the brightness of the image. $B = 0$ returns a black image, $B = 1$ returns the original image.	B	[0, 1]
Color	Adjusts the color balance of the image like in a TV. $C = 0$ returns a black & white image, $C = 1$ returns the original image.	C	[0, 1]
Contrast	Controls the contrast of the image. A $C = 0$ returns a gray image, $C = 1$ returns the original image.	C	[0, 1]
Cutout	Sets a random square patch of side-length ($L \times$ image width) pixels to gray.	L	[0, 0.5]
Equalize	Equalizes the image histogram, and then blends with the original image with blending ratio λ .	λ	[0, 1]
Invert	Inverts the pixels of the image, and then blends with the original image with blending ratio λ .	λ	[0, 1]
Identity	Returns the original image.		
Posterize	Reduces each pixel to B bits.	B	[1, 8]
Rescale	Takes a center crop that is of side-length ($L \times$ image width), and rescales to the original image size using method M .	L M	[0.5, 1.0] see caption
Rotate	Rotates the image by θ degrees.	θ	[-45, 45]
Sharpness	Adjusts the sharpness of the image, where $S = 0$ returns a blurred image, and $S = 1$ returns the original image.	S	[0, 1]
Shear_x	Shears the image along the horizontal axis with rate R .	R	[-0.3, 0.3]
Shear_y	Shears the image along the vertical axis with rate R .	R	[-0.3, 0.3]
Smooth	Adjusts the smoothness of the image, where $S = 0$ returns a maximally smooth image, and $S = 1$ returns the original image.	S	[0, 1]
Solarize	Inverts all pixels above a threshold value of T .	T	[0, 1]
Translate_x	Translates the image horizontally by ($\lambda \times$ image width) pixels.	λ	[-0.3, 0.3]
Translate_y	Translates the image vertically by ($\lambda \times$ image height) pixels.	λ	[-0.3, 0.3]

Experimento: Augmentation e Cutout

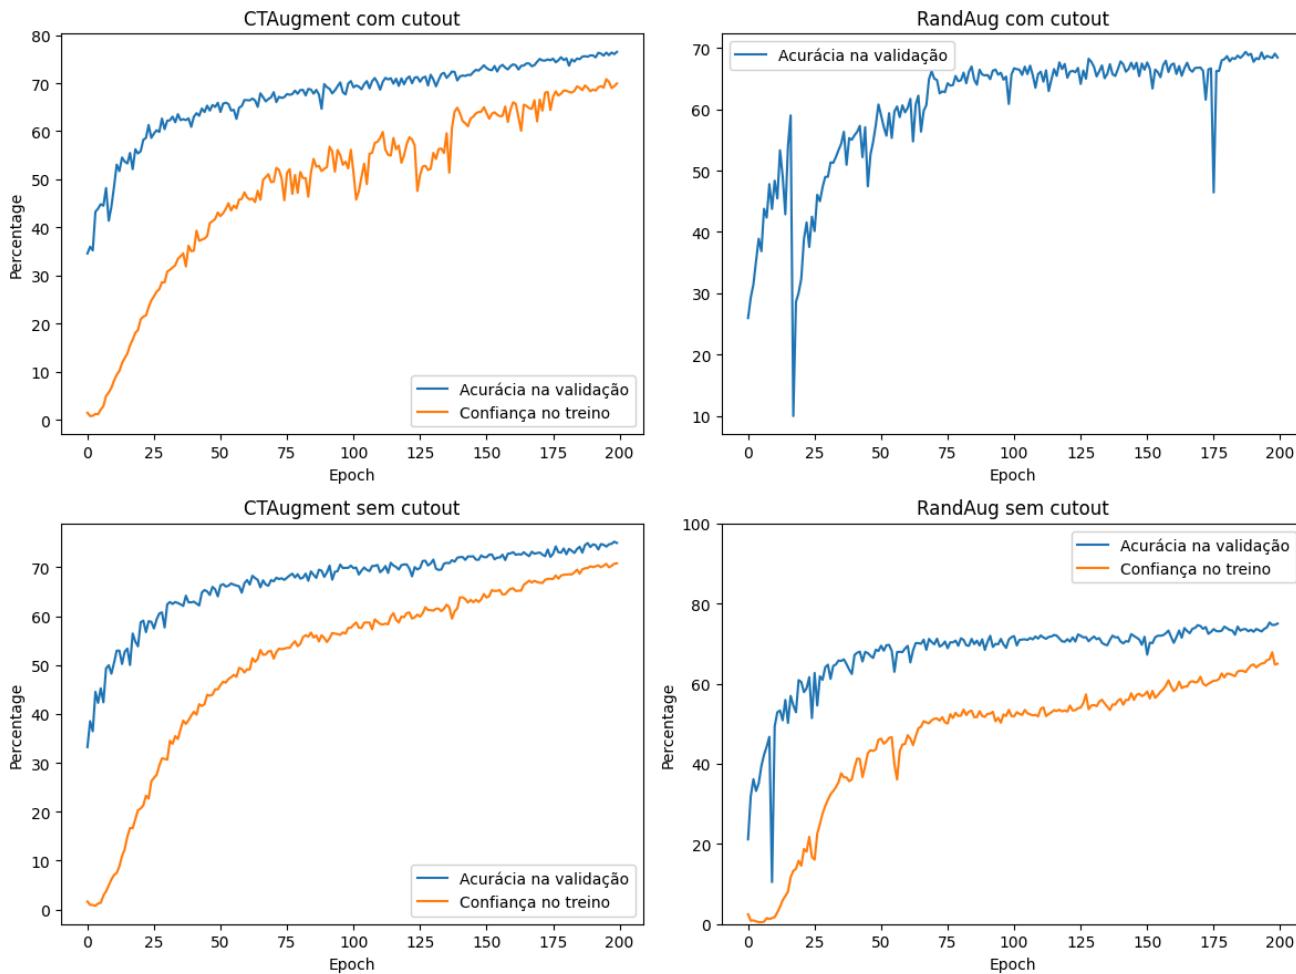
Experimento: Augmentation e Cutout

22/29



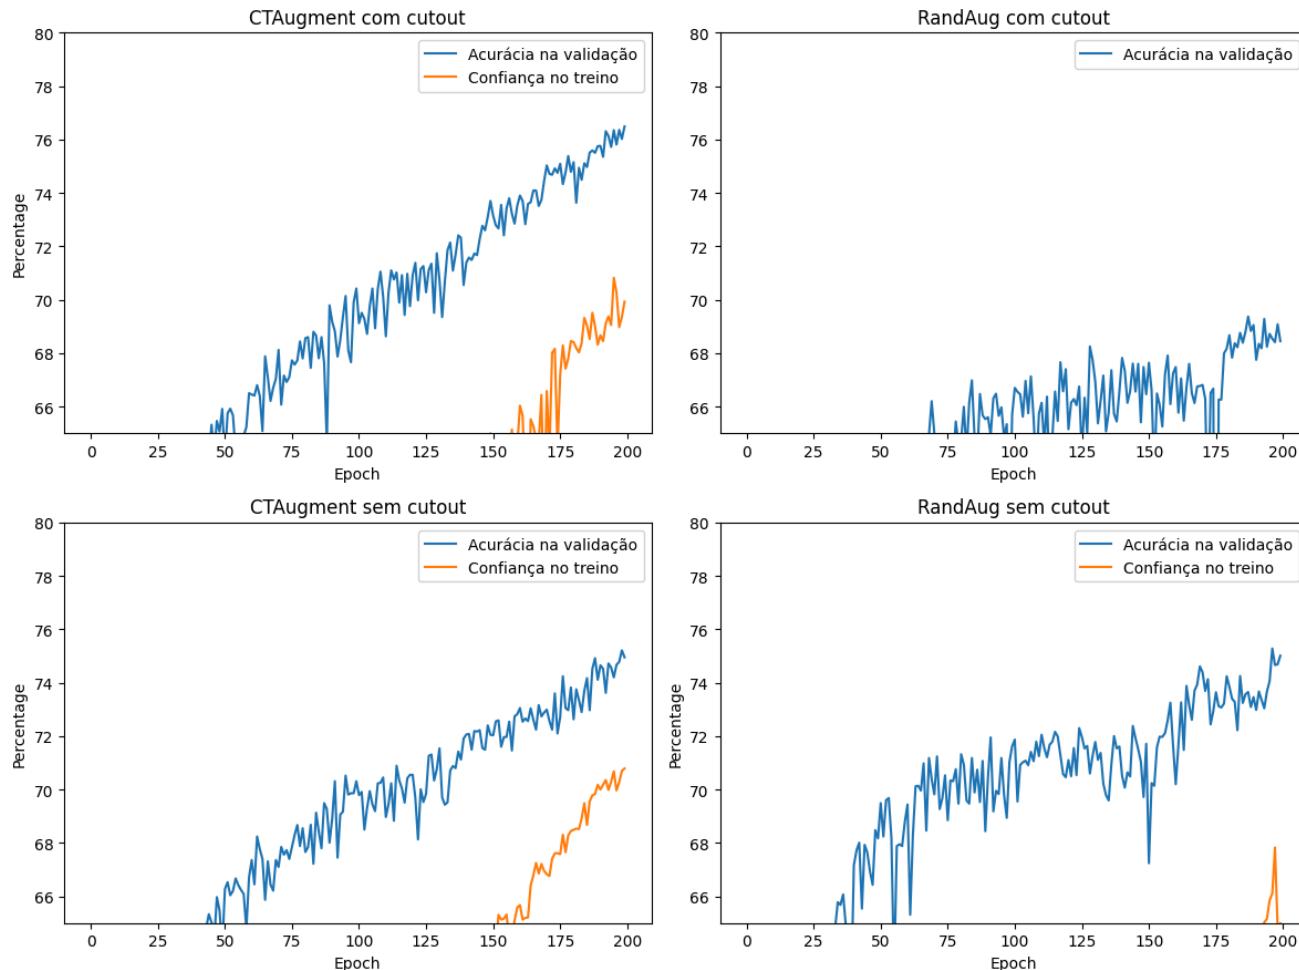
Experimento: Augmentation e Cutout

23/29



Experimento: Augmentation e Cutout

24/29



Experimento: Imagens pouco representativas

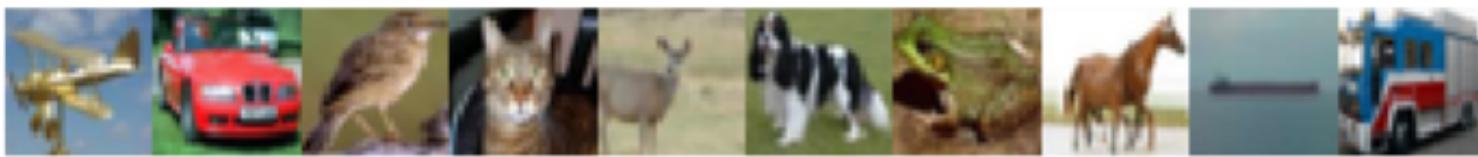


Figure 2: FixMatch reaches 78% CIFAR-10 accuracy using only above 10 labeled images.

```
class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, padding=1)
        self.conv3 = nn.Conv2d(64, 64, 3, padding=1)

        self.fc1 = nn.Linear(64 * 8 * 8, 64)
        self.fc2 = nn.Linear(64, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = F.relu(self.conv3(x))
        x = torch.flatten(x, 1)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

Verificando Influencia da representatividade

28/29

Labeled Dataset Images: Least Representatives



Labeled Dataset Images: Most Representatives



